

Lab9

2022-11-02

Tasks

- ▶ Loops: `for()`, `while()`
- ▶ Loop-related: `apply`, `lapply`, `sapply`, `replicate`
- ▶ Sampling in R: `sample`, `sample_n`, `rep_sample_n`
- ▶ Use one or more of these tools to solve Exercise 3 (Lab 6)
- ▶ More on conditional statements
- ▶ Lab 9 Exercise 1 and Lab 9 Exercise 2
- ▶ R Quiz 9: two chapters from Intermediate R on DataCamp (“Conditionals and Control Flow” and “Loops”)

Conditional statements

The main conditional statements in R are:

- ▶ `if`: we saw this in Lab 8
- ▶ `if; else`: we saw this in Lab 8
- ▶ `else if`: new!

Conditional statements: else if

Basic syntax

```
if(condition){ #if condition is true  
    statement #do something  
}else if(condition 2){  
    statement 2 #do this other thing  
}else{ #if neither statement nor statement 2 are true  
    statement 3 # do this other other thing  
}
```

Conditional statements: else if

```
lunch<-"Dog food"  
if(lunch=="Pizza"){ #if first statatement  
  print("Yay!")  
}else if(lunch=="Rice & beans"){  
  print("Sounds good!")  
}else{  
  print("Meh")  
}
```

```
## [1] "Meh"
```

Conditional statements: else if

```
#you can use negatives too
lunch<-"Dog food"
if(lunch!="Pizza" & lunch!="Rice & beans"){ #if first state
  print("Meh")
}else if(lunch=="Rice & beans"){
  print("Sounds good!")
}else{
  print("Yay!")
}
```

```
## [1] "Meh"
```

Loops in R

Three kinds of loops:

- ▶ for loop
- ▶ while loop
- ▶ repeat loop [we won't discuss]

Loops in R: for loops

Basic syntax

```
for (value in sequence){  
  statement #DO SOMETHING  
}
```


Loops in R: for loops

```
# iterate over elements of a sequence  
for (i in 1: 4){  
  
    print(i * 2)  
  
}
```

```
## [1] 2  
## [1] 4  
## [1] 6  
## [1] 8
```

Loops in R: for loops

```
#iterate over elements of a vector  
for (i in c(-8^2, 2*4, 239, 29321)){  
  
  print(i)  
  
}
```

```
## [1] -64
```

```
## [1] 8
```

```
## [1] 239
```

```
## [1] 29321
```

Loops in R: for loops

#create vector outside of loop and then use it to iterate

```
x <- c(-8^2, 2*4, 239, 29321)
```

```
for (i in x){
```

```
  print(i)
```

```
}
```

```
## [1] -64
```

```
## [1] 8
```

```
## [1] 239
```

```
## [1] 29321
```

Loops in R: while loops

Basic syntax

```
while ( condition ) {  
  statement #do something  
}
```

Loops in R: while loops

```
y<-0
while (y <= 5){ #logical
  # statements
  print(y) # if y <=5, print it
  y = y + 1 #if y <=5 add 1 to it
}
```

```
## [1] 0
```

```
## [1] 1
```

```
## [1] 2
```

```
## [1] 3
```

```
## [1] 4
```

```
## [1] 5
```

```
y
```

```
## [1] 6
```

Loops in R: while loops

```
# whose factorial will be calculated
n <- 5
factorial <- 1 # assigning the factorial variable
i <- 1 # and iteration variable to 1
# using while loop
while (i <= n){
  # multiplying the factorial variable
  # with the iteration variable
  factorial = factorial * i
  # incrementing the iteration variable
  i = i + 1
}
print(factorial) # displaying the factorial

## [1] 120
```

Looping in the command line

`lapply()`: Loop over a list and evaluate a function on each element

`sapply()`: Same as `lapply` but try to simplify the result

`apply()`: Apply a function over the margins of an array

There are others, like `tapply()`, `mapply()`, but we won't discuss them.

Looping in the command line: `lapply()`

- ▶ it loops over a list, iterating over each element in that list
- ▶ it applies a function to each element of the list (a function that you specify)
- ▶ and returns a list (the `l` is for “list”).

Looping in the command line: `lapply()`

Basic syntax for `lapply` and `sapply`

```
lapply(X,FUN)
```

```
lapply(Data, Function to apply to data)
```

Looping in the command line: lapply()

```
#rnorm samples numbers randomly from a normal distribution  
set.seed(123)  
x <- list(a = 1:4, b = rnorm(10), c = rnorm(20, 1), d = rnorm(100, 5))  
#str(x)  
lapply(x, mean)
```

```
## $a  
## [1] 2.5  
##  
## $b  
## [1] 0.07462564  
##  
## $c  
## [1] 0.8920315  
##  
## $d  
## [1] 5.021617
```

Looping in the command line:sapply()

```
#rnorm samples numbers randomly from a normal distribution  
set.seed(123)  
x <- list(a = 1:4, b = rnorm(10), c = rnorm(20, 1), d = rnorm(30, 2))  
#str(x)  
sapply(x, mean)
```

```
##           a           b           c           d  
## 2.50000000 0.07462564 0.89203154 5.02161711
```

Looping in the command line: `apply()`

Basic syntax

```
apply(X,MARGIN, FUN)
```

```
lapply(Data, Where the Function will be applied, Function to apply)
```

```
#Margin: 1 is rows
```

```
#Margin: 2 is cols
```

Looping in the command line: `apply()`

```
set.seed(123)
#make a matrix
mat<-matrix(rnorm(20,1), ncol=2)
apply(mat,1, mean) # 1 signifies rows
```

```
## [1] 1.3318031 1.0648182 1.9797399 1.0905956 0.7867
## [7] 1.4793833 -0.6158392 1.0072515 0.5407733
```

```
rowMeans(mat) #equivalent
```

```
## [1] 1.3318031 1.0648182 1.9797399 1.0905956 0.7867
## [7] 1.4793833 -0.6158392 1.0072515 0.5407733
```

Looping in the command line: `apply()`

```
set.seed(123)
#make a matrix
mat<-matrix(rnorm(20,1), ncol=2)
apply(mat,2, mean) # signifies columns
```

```
## [1] 1.074626 1.208622
```

```
colMeans(mat) #equivalent
```

```
## [1] 1.074626 1.208622
```

The replicate function

```
#replicate the value 3 exactly 10 times  
replicate(n=10, 3)
```

```
## [1] 3 3 3 3 3 3 3 3 3 3
```

```
#replicate the letter 'A' exactly 7 times  
replicate(n=7, 'A')
```

```
## [1] "A" "A" "A" "A" "A" "A" "A"
```

The replicate function

```
set.seed(1)  
#generate 3 values that follow normal distribution (replicate)  
replicate(n=4, rnorm(3, mean=0, sd=1))
```

```
##           [,1]      [,2]      [,3]      [,4]  
## [1,] -0.6264538  1.5952808  0.4874291 -0.3053884  
## [2,]  0.1836433  0.3295078  0.7383247  1.5117812  
## [3,] -0.8356286 -0.8204684  0.5757814  0.3898432
```


Practice!

Lab9_ex2.Rmd

Sampling in R

`sample(x = , size = , replace = , prob =)`: Generate a sample of size `size`, from a vector `x`, with (`replace = TRUE`) or without (`replacement = FALSE`) replacement. By default the size is the length of `x`, sampling occurs without replacement and probabilities are equal. Change these defaults by specifying a value for the argument. For example, to have unequal sampling probabilities, include a vector of length `x`, in which the i^{th} entry describes the relative probability of sampling the i^{th} value in `x`.

`sample_n(tbl = , size = , replace = , weight =)`: Generate a sample of size `size`, from a tibble `tbl`, with (`replace = TRUE`) or without (`replacement = FALSE`) replacement. All arguments are the same as in `sample()` except `weight` replaces `prob`, and `tbl` replaces `x`. `sample_n()` is a function in the `dplyr` package, which is loaded with `tidyverse`.

New functions for sampling in R!

These functions extend the functionality of `dplyr::sample_n()` and `dplyr::slice_sample()` by allowing for repeated sampling of data. This operation is especially helpful while creating sampling distributions

These are related to `dplyr::sample_n` and `'dplyr::slice_sample`, respectively:

`rep_sample_n()`: Generate `reps` samples of size `size`

`rep_sample_slice()`: Generate `reps` samples of size `n`

Both return a tibble grouped by replicate. ‘

New functions for sampling in R!

```
library(infer) # a tidymodels package  
#sample 2 rows from the iris dataset without replacement.  
set.seed(123)  
test1<-iris %>%  
  rep_sample_n(size = 3, replace = F, reps = 2) #uses size  
test1
```

```
## # A tibble: 6 x 6  
## # Groups:   replicate [2]  
##   replicate Sepal.Length Sepal.Width Petal.Length Petal.  
##           <int>         <dbl>         <dbl>         <dbl>  
## 1           1           4.3           3           1.1  
## 2           1           5           3.3           1.4  
## 3           1           7.7           3.8           6.7  
## 4           2           4.4           3.2           1.3  
## 5           2           4.3           3           1.1  
## 6           2           7.7           3.8           6.7
```

New functions for sampling in R!

```
library(infer) # a tidymodels package  
#sample 2 rows from the iris dataset without replacement.  
set.seed(123)  
test2<-iris %>%  
  rep_slice_sample(n = 3, replace = F, reps = 2) #uses n in  
test2
```

```
## # A tibble: 6 x 6
```

```
## # Groups:   replicate [2]
```

```
##   replicate Sepal.Length Sepal.Width Petal.Length Petal.
```

```
##           <int>         <dbl>         <dbl>         <dbl>
```

```
## 1             1           4.3           3           1.1
```

```
## 2             1           5           3.3           1.4
```

```
## 3             1           7.7           3.8           6.7
```

```
## 4             2           4.4           3.2           1.3
```

```
## 5             2           4.3           3           1.1
```

```
## 6             2           7.7           3.8           6.7
```

New functions for sampling in R!

```
test1==test2
```

##	replicate	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
## [1,]	TRUE	TRUE	TRUE	TRUE	TRUE
## [2,]	TRUE	TRUE	TRUE	TRUE	TRUE
## [3,]	TRUE	TRUE	TRUE	TRUE	TRUE
## [4,]	TRUE	TRUE	TRUE	TRUE	TRUE
## [5,]	TRUE	TRUE	TRUE	TRUE	TRUE
## [6,]	TRUE	TRUE	TRUE	TRUE	TRUE

Back to the Exercise: Making a distribution of sample means (from lab 6)

Last week you learned how to take a sample (random rows) from a dataset and calculate statistics on it.

Specifically, you took two independent samples from the `human_genes` dataset and calculated some descriptive statistics.

Exercise: Making a distribution of sample means

Step 1: Read in the `human_genes.csv` dataset and get summaries

```
library(dplyr) #load the dplyr package
library(readr) #load the readr package
#read in the human genes dataset
human_genes<-readr::read_csv("input_files/human_genes.csv")
```

```
## Rows: 22385 Columns: 4
## -- Column specification -----
## Delimiter: ","
## chr (3): gene, name, description
## dbl (1): size
##
## i Use `spec()` to retrieve the full column specification
## i Specify the column types or set `show_col_types = FALSE`
```


Exercise: Making a distribution of sample means

Step 1: Read in the `human_genes.csv` dataset and get summaries

```
glimpse(human_genes) #glimpse the dataset
```

```
## Rows: 22,385
```

```
## Columns: 4
```

```
## $ gene      <chr> "ENSG00000000003.14", "ENSG00000000003.14"
```

```
## $ size      <dbl> 3796, 1339, 1161, 6364, 4355, 2729,
```

```
## $ name      <chr> "TSPAN6", "TNMD", "DPM1", "SCYL3", "
```

```
## $ description <chr> "tetraspanin", "tenomodulin", "dolic
```

```
#remove useless columns
```

```
human_genes<-human_genes %>%
```

```
select(name, size)
```

Exercise: Making a distribution of sample means

Step 2: Calculate the following descriptive statistics for the dataset:
mean, median, standard deviation, mode, IQR.

```
#summarise the human genes dataset creating three columns
human_genes_summ <- human_genes %>%
  summarise(
    MeanLength = mean(size),
    MedianLength = median(size),
    SDLength = sd(size),
    IQR = IQR(size)
  )
human_genes_summ #
```

```
## # A tibble: 1 x 4
##   MeanLength MedianLength SDLength   IQR
##       <dbl>         <dbl>    <dbl> <dbl>
## 1      3511.         2744      2833.  2827
```

Exercise: Making a distribution of sample means

Step 3: take 10,000 random samples of length $n=100$ from the `human_genes` dataset and calculate the same summaries as above.

Wow! Let's break this down:

- ▶ sample 10 random samples of length $n=100$
- ▶ calculate the same summaries as above for each sample
- ▶ if everything works, switch to 10,000

Exercise: Making a distribution of sample means

Step 3a: sample 10 random samples of length $n=100$ (using for)

```
library(dplyr)
x<-seq(1:10) #create a vector with numbers from 1 to 10000
# for each iteration, take a sample of size 100 without rep
samps<-list() #create empty list to put resamps in
for(i in x){ #for loop
  samps[[i]]<-human_genes %>%
    sample_n(size=100, replace=F)
  print(i) #print each iteration
}
```

```
## [1] 1
```

```
## [1] 2
```

```
## [1] 3
```

```
## [1] 4
```

```
## [1] 5
```

```
## [1] 6
```

```
## [1] 7
```

Exercise: Making a distribution of sample means

Step 3b: calculate the same summaries as above for each sample

```
res<-list() #create empty list to put the summaries in
for(i in 1:length(samps)){ #for loop for each elements in
  res[[i]]<- samps[[i]] %>%
    summarise(
      MeanLength = mean(size),
      MedianLength = median(size),
      SDLenght = sd(size),
      IQR = IQR(size)
    )
  print(i) #print each iteration
}
```

```
## [1] 1
```

```
## [1] 2
```

```
## [1] 3
```

```
## [1] 4
```

```
## [1] 5
```

Exercise: Making a distribution of sample means

Can we do it all in one go?

```
x<-seq(1:10) #create a vector with numbers from 1 to 10000
#create tibble with the columns produced by our summaries
#dummy tibble
resamps<-tibble(MeanLength=NA, MedianLength=NA, SLength=NA)
for(i in x){ #for loop
  temp<-human_genes %>% #assign to temp object
  sample_n(size=100, replace=F) %>% #sample
  summarise( # summaries
    MeanLength = mean(size),
    MedianLength = median(size),
    SLength = sd(size),
    IQR = IQR(size)
  )
  resamps<-bind_rows(resamps, temp) #add new row to tibble
  print(i) #print each iteration
}
```

Exercise: Making a distribution of sample means

#the first row has NAs. we can easily get rid of it.

```
resamps<-na.omit(resamps) #remove lines with NA  
resamps
```

```
## # A tibble: 10 x 4
```

```
##      MeanLength MedianLength SDLength    IQR  
##      <dbl>          <dbl>    <dbl> <dbl>  
##  1      3269.          2716.    2187. 2952.  
##  2      3772.          3020.    2489. 3320.  
##  3      3704.          2936     3283. 2936.  
##  4      3945.          3260     2664. 2564.  
##  5      3506.          2828     2611. 3014.  
##  6      4064.          2756.    4847. 3218.  
##  7      3220.          2633     2404. 2863  
##  8      3362.          2709     2550. 2629.  
##  9      3377.          2869     2132. 3260.  
## 10      3437.          2380     3193. 2570.
```

Exercise: Making a distribution of sample means

Step 3c: if everything works, switch to 10,000

```
nsamps=10000
resamps<-tibble(MeanLength=NA, MedianLength=NA, SDLength=NA)
for(i in 1:nsamps){ #for loop
  temp<-human_genes %>% #assign to temp object
    sample_n(size=100, replace=F) %>% #sample
    summarise( # summaries
      MeanLength = mean(size),
      MedianLength = median(size),
      SDLength = sd(size),
      IQR = IQR(size)
    )
  resamps<-bind_rows(resamps, temp) #add new row to tibble
  if(i %in% seq(from=1, to=nsamps, by=500)){
    print(i) #print 1st, 501st, 1001st etc iterations
  }
}
```


Exercise: Making a distribution of sample means

Another option: using `rep_slice_sample`

Make a function that calculates all the summaries


```
library(infer)
#samps
nsamps<-10 #once this works, switch to 10,000
reps<-rep_sample_n(human_genes, size=100, replace=FALSE, re
  group_by(replicate) %>%
  summarise( # summaries
    MeanLength = mean(size),
    MedianLength = median(size),
    SDLength = sd(size),
    IQR = IQR(size)
  ) %>%
  ungroup
```

Exercise: Making a distribution of sample means

Step 4: Plot the sampling distribution with ggplot

```
library(ggplot2)
#rerun previous code with nsamps=10000
ggplot(reps, aes(x=MeanLength)) +
  geom_histogram() +
  theme_minimal()
```

`stat_bin()` using `bins = 30`. Pick better value with `



Exercise: Making a distribution of sample means

Final step: Obtain the grand mean (mean of sample means) and the standard deviation of the sample means and: compare with the population mean; calculate SEM using the distribution and the formula and compare.

#Try it!

Exercises!

- ▶ Binomial hypothesis testing: `lab9_ex1.Rmd`
- ▶ Loops & Conditional Statements: Exercises in `lab9_ex2.Rmd`.